

Question_3

November 4, 2020

1 CS5590: Foundations of Machine Learning

1.1 Assignment 1

1.2 Question 3

1.3 Authors

Name	Roll number
Raj Patil	CS18BTECH11039
Karan Bhukar	CS18BTECH11021

2 Question 3

```
[5]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import io
import matplotlib.pyplot as plt
from matplotlib.pyplot import figure
from scipy.special import factorial
from mpl_toolkits.axes_grid1 import host_subplot
import mpl_toolkits.axisartist as AA
```

```
[6]: data = pd.read_csv('HorseKicks.txt',delimiter='\t',index_col=0).transpose()
data['Totals']=data.sum(axis=1)
d=data;d
```

```
[6]: Year  1875  1876  1877  1878  1879  1880  1881  1882  1883  1884  ...  1886  \
GC        0    2    2    1    0    0    1    1    0    3  ...    2
C1        0    0    0    2    0    3    0    2    0    0  ...    1
C2        0    0    0    2    0    2    0    0    1    1  ...    0
C3        0    0    0    1    1    1    2    0    2    0  ...    0
C4        0    1    0    1    1    1    1    0    0    0  ...    1
C5        0    0    0    0    2    1    0    0    1    0  ...    1
C6        0    0    1    0    2    0    0    1    2    0  ...    1
C7        1    0    1    0    0    0    1    0    1    1  ...    0
```

C8	1	0	0	0	1	0	0	1	0	0	...	0
C9	0	0	0	0	0	2	1	1	1	0	...	1
C10	0	0	1	1	0	1	0	2	0	2	...	0
C11	0	0	0	0	2	4	0	1	3	0	...	1
C14	1	1	2	1	1	3	0	4	0	1	...	3
C15	0	1	0	0	0	0	0	1	0	1	...	0

Year	1887	1888	1889	1890	1891	1892	1893	1894	Totals
GC	1	0	0	1	0	1	0	1	16
C1	1	1	0	2	0	3	1	0	16
C2	2	1	1	0	0	2	0	0	12
C3	1	0	1	2	1	0	0	0	12
C4	0	0	0	0	1	1	0	0	8
C5	0	1	1	1	1	1	1	0	11
C6	3	1	1	1	0	3	0	0	17
C7	2	0	0	2	1	0	2	0	12
C8	1	0	0	0	1	1	0	1	7
C9	1	0	1	2	0	1	0	0	13
C10	0	0	2	1	3	0	1	1	15
C11	1	1	2	1	3	1	3	1	25
C14	2	1	0	2	1	1	0	0	24
C15	0	0	2	2	0	0	0	0	8

[14 rows x 21 columns]

```
[7]: train = d.iloc[:, :13]
test = d.iloc[:, -8:-1]
```

```
[8]: train.head()
```

```
[8]: Year  1875  1876  1877  1878  1879  1880  1881  1882  1883  1884  1885  1886  \
GC        0    2    2    1    0    0    1    1    0    3    0    2
C1        0    0    0    2    0    3    0    2    0    0    0    1
C2        0    0    0    2    0    2    0    0    1    1    0    0
C3        0    0    0    1    1    1    2    0    2    0    0    0
C4        0    1    0    1    1    1    1    0    0    0    0    1
```

```
Year  1887
GC    1
C1    1
C2    2
C3    1
C4    0
```

```
[9]: test.head()
```

```
[9]: Year  1888  1889  1890  1891  1892  1893  1894
      GC      0    0    1    0    1    0    1
      C1      1    0    2    0    3    1    0
      C2      1    1    0    0    2    0    0
      C3      0    1    2    1    0    0    0
      C4      0    0    0    1    1    0    0
```

```
[10]: test.sum(axis=1)
```

```
[10]: GC      3
      C1      7
      C2      4
      C3      4
      C4      2
      C5      6
      C6      6
      C7      5
      C8      3
      C9      4
      C10     8
      C11     12
      C14      5
      C15      4
      dtype: int64
```

3 Description

- modelling the data as a poisson distribution, independently for each corp

$$N_{deaths} \sim \mathcal{P}(\theta)$$

- note that there are no input features in this case and only a parameter dependent output
- once we find θ from the training set, the prediction will be the mode of the distribution i.e $\lfloor \theta \rfloor$ if we are predicting the answer for one year
- but in this case(seven years), the prediction will be $\lfloor 7(\theta) \rfloor$ as θ represents the death rate per year

3.1 Question 3.1:

3.1.1 ML estimate

- for a corp, the number of deaths every year independently and identically follow our poisson model

- hence for a corp, the likelihood over D years (with death d_i for year i) is given by...

$$\mathcal{L}(\theta) = \prod_{i=1}^D \frac{\theta^{d_i} e^{-\theta}}{(d_i)!} \quad (1)$$

$$\therefore \ln(\mathcal{L}(\theta)) = \sum_{i=1}^D (d_i \ln(\theta) - \theta - \ln((d_i)!)) \quad (2)$$

$$\theta_{ML} = \underset{\theta}{argmax} (\ln(\mathcal{L}(\theta))) \quad (3)$$

$$\text{which occurs at} \quad (4)$$

$$\nabla_{\theta} (\ln(\mathcal{L}(\theta))) = 0 = \sum_{i=1}^D \left(\frac{d_i}{\theta} - 1 \right) \quad (5)$$

$$\therefore \theta_{ML} = \frac{\sum_{i=1}^D d_i}{D} \quad (6)$$

- this is simply the mean of the deaths over the years
- note that we do not have any features to account for and we are estimating θ directly hence we have a closed form solution

```
[11]: theta_ML = train.mean(axis=1)
      theta_ML
```

```
[11]: GC      1.000000
      C1      0.692308
      C2      0.615385
      C3      0.615385
      C4      0.461538
      C5      0.384615
      C6      0.846154
      C7      0.538462
      C8      0.307692
      C9      0.692308
      C10     0.538462
      C11     1.000000
      C14     1.461538
      C15     0.307692
      dtype: float64
```

4 predictions

```
[12]: years = test.shape[1]
      predictions = pd.Series(np.floor(theta_ML*years), dtype=np.int)
      predictions
```

```
[12]: GC      7
      C1      4
      C2      4
      C3      4
      C4      3
      C5      2
      C6      5
      C7      3
      C8      2
      C9      4
      C10     3
      C11     7
      C14    10
      C15     2
      dtype: int64
```

```
[13]: actuals = test.sum(axis=1)
      actuals
```

```
[13]: GC      3
      C1      7
      C2      4
      C3      4
      C4      2
      C5      6
      C6      6
      C7      5
      C8      3
      C9      4
      C10     8
      C11    12
      C14     5
      C15     4
      dtype: int64
```

5 Calculating error

$$E_{RMS} = \sqrt{\frac{\sum_{C_i \in C} (y_{C_i} - \hat{y}_{C_i})^2}{|C|}}$$

```
[14]: rmse_ML = np.sqrt(
      (
          ((actuals - predictions)**2).sum()
          /
          years
```

```
)  
)  
rmse_ML
```

[14]: 4.2594432902501635

5.1 Question 3.2

5.2 MAP estimate

Assuming the prior to be an exponential distribution by the following intuition: - the chances of lower rate of deaths are significantly greater than that of a higher rate of deaths

$$\theta \sim \mathcal{E}(\lambda)$$

$$f(\theta) = \lambda e^{-\lambda\theta}$$

choosing the value of λ to be 1

$$f(\theta) = e^{-\theta}$$

- Why not generalized gamma?:
 - although exponential is a subset of the gamma function, we don't prefer to use the general gamma as that involves deliberately choosing a finite non-zero peak which will involve some sort of a look at the number of deaths from the training set
 - one could say that the mode of the prior is being selected as a hyper-parameter and then one could distribute the training set further into a validation set for selecting the hyper-parameter but that defeats the notion of a “prior”

5.3 Maximizing the posterior...

... after viewing data for D years

$$p(\theta|D) = \frac{p(D|\theta) \cdot p(\theta)}{p(D)}$$

here $p(D|\theta)$ is the likelihood as given before :

$$p(D|\theta) = \mathcal{L}(\theta) = \prod_{i=1}^D \frac{\theta^{d_i} e^{-\theta}}{(d_i)!}$$

Maximizing the natural log of the posterior will yield the same results due the monotonic nature of the logarithmic function

$$\therefore \ln(p(\theta|D)) = \ln(p(D|\theta)) + \ln(p(\theta)) - \ln(p(D))$$

$$\theta_{MAP} = \underset{\theta}{\operatorname{argmax}} \ln(p(\theta|D))$$

$$\nabla_{\theta} \ln(p(D|\theta)) + \ln(p(\theta)) = 0$$

$$\nabla_{\theta} \sum_{i=1}^D (d_i \cdot \ln(\theta) - \theta) - \theta = 0$$

$$\theta_{MAP} = \frac{\sum_{i=1}^D d_i}{D+1}$$

```
[15]: train
```

```
[15]: Year  1875  1876  1877  1878  1879  1880  1881  1882  1883  1884  1885  1886  \
GC         0     2     2     1     0     0     1     1     0     3     0     2
C1         0     0     0     2     0     3     0     2     0     0     0     1
C2         0     0     0     2     0     2     0     0     1     1     0     0
C3         0     0     0     1     1     1     2     0     2     0     0     0
C4         0     1     0     1     1     1     1     0     0     0     0     1
C5         0     0     0     0     2     1     0     0     1     0     0     1
C6         0     0     1     0     2     0     0     1     2     0     1     1
C7         1     0     1     0     0     0     1     0     1     1     0     0
C8         1     0     0     0     1     0     0     1     0     0     0     0
C9         0     0     0     0     0     2     1     1     1     0     2     1
C10        0     0     1     1     0     1     0     2     0     2     0     0
C11        0     0     0     0     2     4     0     1     3     0     1     1
C14        1     1     2     1     1     3     0     4     0     1     0     3
C15        0     1     0     0     0     0     0     1     0     1     1     0

Year  1887
GC     1
C1     1
C2     2
C3     1
C4     0
C5     0
C6     3
C7     2
C8     1
C9     1
C10    0
C11    1
C14    2
C15    0
```

```
[16]: theta_MAP = train.sum(axis=1) / (train.shape[1]+1)
      theta_MAP[2]
```

```
[16]: 0.5714285714285714
```

```
[17]: # sanity check
      theta_MAP < theta_ML
```

```
[17]: GC      True
      C1      True
      C2      True
      C3      True
      C4      True
      C5      True
      C6      True
      C7      True
      C8      True
      C9      True
      C10     True
      C11     True
      C14     True
      C15     True
      dtype: bool
```

6 predictions

```
[18]: years = test.shape[1]
      predictions = pd.Series(np.floor(theta_MAP*years),dtype=np.int)
      predictions
```

```
[18]: GC      6
      C1      4
      C2      4
      C3      4
      C4      3
      C5      2
      C6      5
      C7      3
      C8      2
      C9      4
      C10     3
      C11     6
      C14     9
      C15     2
      dtype: int64
```

```
[19]: actuals = test.sum(axis=1)
      actuals
```

```
[19]: GC      3
      C1      7
      C2      4
      C3      4
      C4      2
```



```

C5      6
C6      6
C7      5
C8      3
C9      4
C10     8
C11    12
C14     5
C15     4
dtype: int64

```

7 Calculating error

$$E_{RMS} = \sqrt{\frac{\sum_{C_i \in C} (y_{C_i} - \hat{y}_{C_i})^2}{|C|}}$$

```

[20]: rmse_MAP = np.sqrt(
      (
          ((actuals - predictions)**2).sum()
          /
          years
      )
      )
rmse_MAP

```

```
[20]: 4.174754056057845
```

```
[21]: rmse_MAP < rmse_ML
```

```
[21]: True
```

8 This solidifies that our intuition is somewhat sane

8.1 Plots

```

[22]: def likelihood(theta, total_deaths, pro_fac):
      return (pow(theta, total_deaths) * (np.exp(theta * (-13)))) / pro_fac

```

```

[23]: def prior(x):
      return np.exp(-1 * x)

```

```

[24]: def posterior(theta, total_deaths, pro_fac):
      return (pow(theta, total_deaths) * (np.exp(theta * (-13)))) * (np.exp(-1 * theta)) /
      ↪ (pro_fac)

```

```
[25]: def plot(c_num, train, theta_MAP, theta_ML):
    total_deaths = train.sum(axis=1)[c_num]
    print(total_deaths)
    pro_fac = 0
    for i in range(0,13):
        pro_fac = pro_fac + factorial(train.iat[c_num,i],exact = True)
    figure(num=None, figsize=(20, 12), dpi=80, facecolor='w', edgecolor='k')
    x = np.arange(0,1.75,.01)
    host = host_subplot(111, axes_class=AA.Axes)
    plt.subplots_adjust(right=0.75)
    par1 = host.twinx()
    par2 = host.twinx()

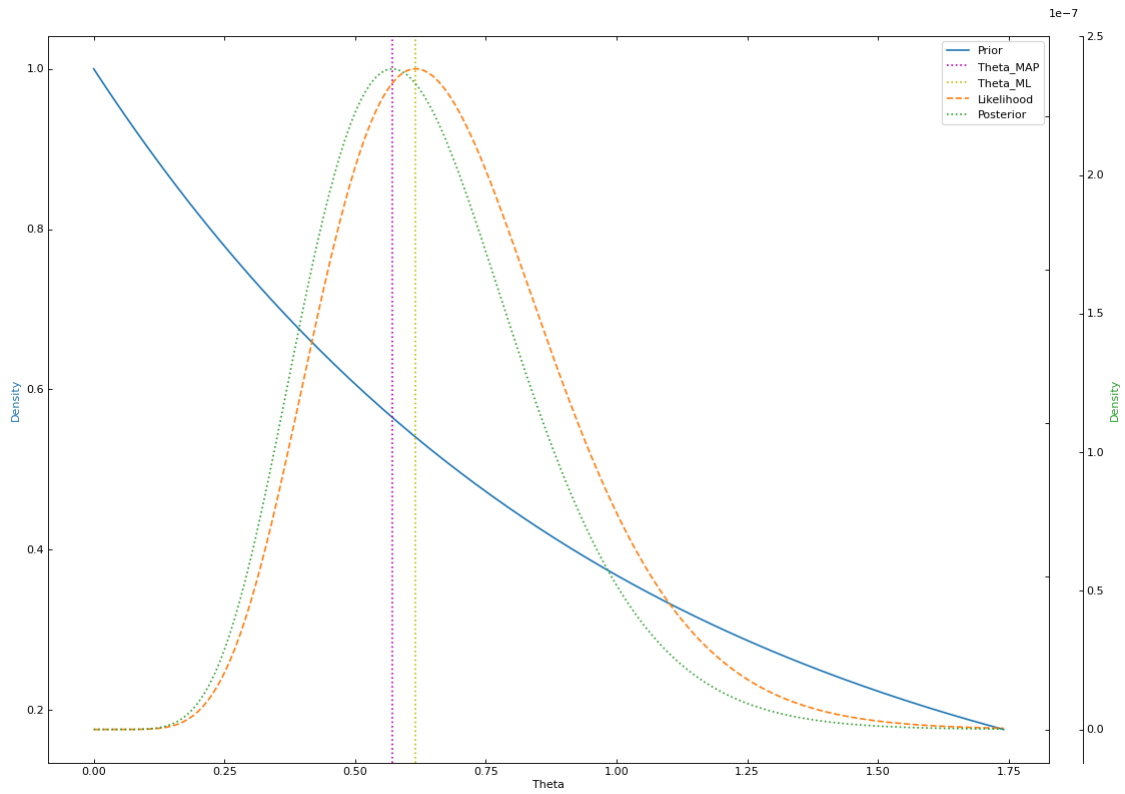
    offset = 30
    new_fixed_axis = par2.get_grid_helper().new_fixed_axis
    par2.axis["right"] = new_fixed_axis(loc="right", axes=par2,
                                       offset=(offset, 0))

    y = np.arange(0,1,.01)

    par2.axis["right"].toggle(all=True)
    host.set_xlabel("Theta")
    host.set_ylabel("Density")
    par1.set_ylabel("Density")
    par2.set_ylabel("Density")
    p1, = host.plot(x,prior(x),label="Prior")
    p2, = par1.
    →plot(x,likelihood(x,total_deaths,pro_fac),linestyle='--',label="Likelihood")
    p3, = par2.plot(x,posterior(x,total_deaths,pro_fac),linestyle=':
    →',label="Posterior")
    plt.axvline(x=theta_MAP,color='m',linestyle="dotted",label="Theta_MAP")
    plt.axvline(x=theta_ML,color='y',linestyle="dotted",label="Theta_ML")
    host.legend()
    host.axis["left"].label.set_color(p1.get_color())
    par1.axis["right"].label.set_color(p2.get_color())
    par2.axis["right"].label.set_color(p3.get_color())
    plt.draw()
    plt.show()
```

8.2 corp 2

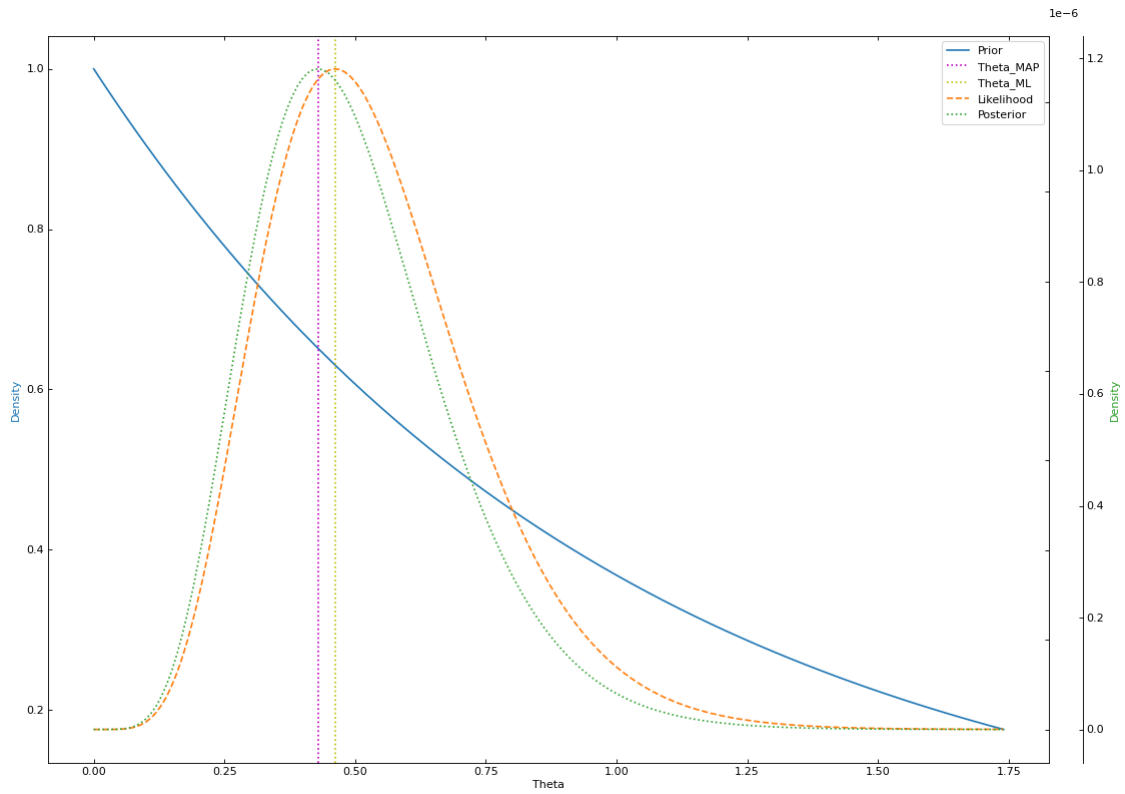
```
[26]: plot(2,train,theta_MAP[2],theta_ML[2])
```



8.3 corp 4

```
[27]: plot(4,train,theta_MAP[4],theta_ML[4])
```

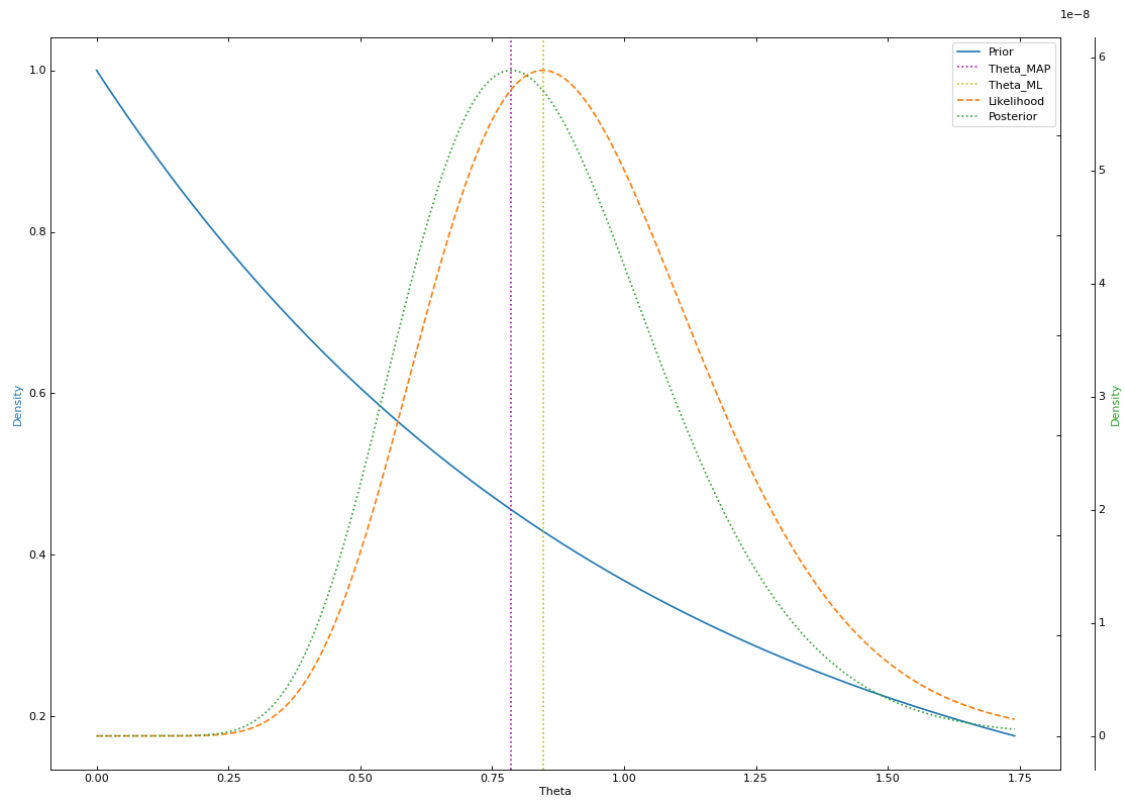
6



8.4 corp 6

```
[28]: plot(6,train,theta_MAP[6],theta_ML[6])
```

11



8.4.1 note that the mode of the posterior lies between that of the prior and the likelihood

as expected