

# Question 4

November 4, 2020

## 1 CS5590: Foundations of Machine Learning

### 1.1 Assignment 1

### 1.2 Question 4

### 1.3 Authors

Name	Roll number
Raj Patil	CS18BTECH11039
Karan Bhukar	CS18BTECH11021

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import random
```

```
[2]: d = pd.read_csv('hour.csv');d.info()
# no null vals
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17379 entries, 0 to 17378
Data columns (total 17 columns):
#   Column          Non-Null Count  Dtype
---  -
0   instant         17379 non-null  int64
1   dteday          17379 non-null  object
2   season          17379 non-null  int64
3   yr              17379 non-null  int64
4   mnth            17379 non-null  int64
5   hr              17379 non-null  int64
6   holiday         17379 non-null  int64
7   weekday         17379 non-null  int64
8   workingday      17379 non-null  int64
9   weathersit       17379 non-null  int64
10  temp            17379 non-null  float64
11  atemp           17379 non-null  float64
12  hum             17379 non-null  float64
```

```

13  windspeed    17379 non-null  float64
14  casual       17379 non-null  int64
15  registered   17379 non-null  int64
16  cnt          17379 non-null  int64
dtypes: float64(4), int64(12), object(1)
memory usage: 2.3+ MB

```

```
[3]: d.iloc[0]
```

```

[3]: instant          1
    dteday      2011-01-01
    season          1
    yr            0
    mnth          1
    hr            0
    holiday        0
    weekday        6
    workingday     0
    weathersit      1
    temp          0.24
    atemp         0.2879
    hum           0.81
    windspeed      0
    casual         3
    registered     13
    cnt           16
    Name: 0, dtype: object

```

### 1.3.1 subquestion 2 and 3

```

[46]: def plot(feature,d):
    if feature == 'season':
        X=["1st","2nd","3rd","4th"]
        Y=[0,0,0,0]
        for i in range(d.iloc[:,2].shape[0]):
            Y[d.iloc[i,2]-1] += 1
        mean = [i/d.iloc[:,2].shape[0] for i in Y]
        for i,me in enumerate(mean):
            print("mean of count for season {} is = {}".format(i+1,me))
        plt.xlabel("Seasons")
        plt.ylabel("Count of Rented Bikes")
        plt.bar(X,Y)

    elif feature == 'yr':
        X=["1st","2nd"]
        Y=[0,0]
        for i in range(d.iloc[:,3].shape[0]):

```

```

        Y[d.iloc[i,3]] += 1
    mean = [i/d.iloc[:,3].shape[0] for i in Y]
    for i,me in enumerate(mean):
        print("mean of count for year {} is = {}".format(i,me))
    plt.xlabel("year")
    plt.ylabel("Count of Rented Bikes")
    plt.bar(X,Y)

elif feature == 'month':
    X=[1,2,3,4,5,6,7,8,9,10,11,12]
    Y=[0,0,0,0,0,0,0,0,0,0,0,0]
    month = ['January', 'February', 'March', 'April', 'May', 'June', 'July', '
    →'August', 'September', 'October', 'November', 'December']
    for i in range(d.iloc[:,4].shape[0]):
        Y[d.iloc[i,4]-1] += 1
    mean = [i/d.iloc[:,4].shape[0] for i in Y]
    for i,me in enumerate(mean):
        print("mean of count for month {} is = {}".format(i+1,me))
    plt.xlabel("Month of the year")
    plt.ylabel("Count of Rented Bikes")
    plt.xticks(X, month, rotation='vertical')
    plt.bar(X,Y)

elif feature == 'holiday':
    X=["No","Yes"]
    Y=[0,0]
    for i in range(d.iloc[:,6].shape[0]):
        Y[d.iloc[i,6]] += 1
    mean = [i/d.iloc[:,6].shape[0] for i in Y]
    for i,me in enumerate(mean):
        print("mean of count for holiday {} is = {}".format(i,me))
    plt.xlabel("Holiday")
    plt.ylabel("Count of Rented Bikes")
    plt.bar(X,Y)


elif feature == 'weekday':
    X=["Monday","Tuesday","Wednesday","Thursday","Friday","Saturday"]
    Y=[0,0,0,0,0,0]
    for i in range(d.iloc[:,7].shape[0]):
        Y[d.iloc[i,7]-1] += 1
    mean = [i/d.iloc[:,7].shape[0] for i in Y]
    for i,me in enumerate(mean):
        print("mean of count for weekday {} is = {}".format(i+1,me))
    plt.xlabel("Day of the week",labelpad=20)
    plt.ylabel("Count of Rented Bikes")
    plt.bar(X,Y)

```

```
plt.show()
```


```
[47]: plot("season",d)
```

```
mean of count for season 1 is = 0.2440876920421198  
mean of count for season 2 is = 0.25369699062086426  
mean of count for season 3 is = 0.25870303239541975  
mean of count for season 4 is = 0.24351228494159619
```




```
[48]: plot("yr",d)
```

```
mean of count for year 0 is = 0.4974394384026699  
mean of count for year 1 is = 0.5025605615973301
```



```
[49]: plot("month",d)
```

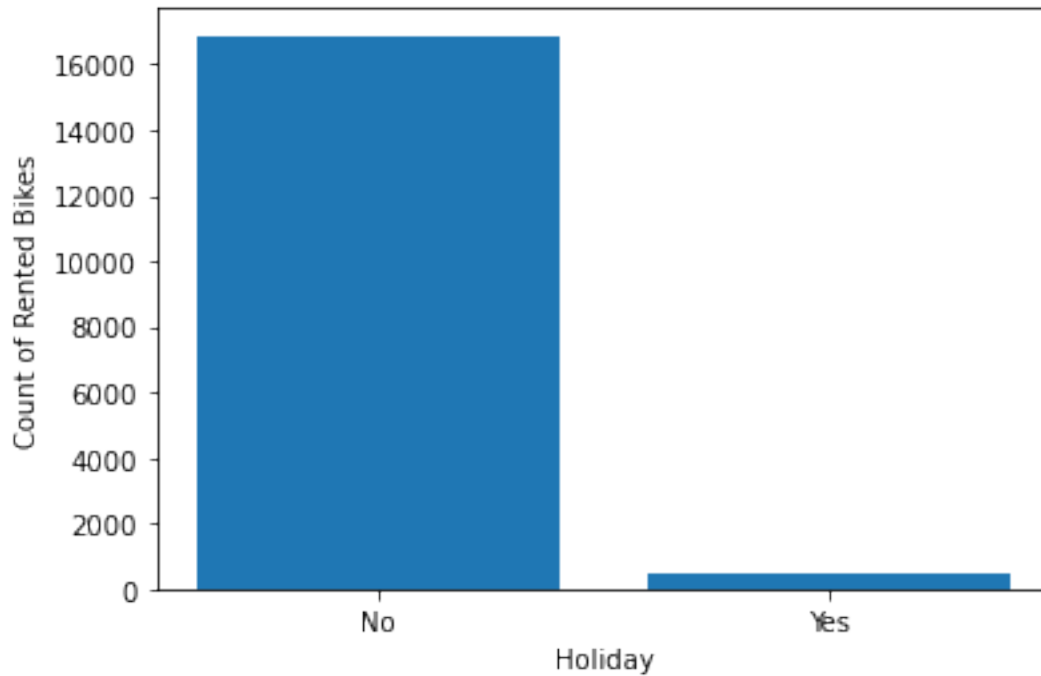
```
mean of count for month 1 is = 0.08222567466482536
mean of count for month 2 is = 0.07716209218021751
mean of count for month 3 is = 0.08475746590712929
mean of count for month 4 is = 0.08268600034524426
mean of count for month 5 is = 0.08562057655791472
mean of count for month 6 is = 0.08285862247540135
mean of count for month 7 is = 0.08562057655791472
mean of count for month 8 is = 0.08487254732723402
mean of count for month 9 is = 0.08268600034524426
mean of count for month 10 is = 0.08349157028597733
mean of count for month 11 is = 0.08268600034524426
mean of count for month 12 is = 0.08533287300765291
```



```
[50]: plot("holiday",d)
```


```
mean of count for holiday 0 is = 0.971229644973819
```

```
mean of count for holiday 1 is = 0.028770355026181024
```



```
[51]: plot("weekday",d)
```

```
mean of count for weekday 1 is = 0.1426434202198055  
mean of count for weekday 2 is = 0.1411473617584441  
mean of count for weekday 3 is = 0.14241325737959606  
mean of count for weekday 4 is = 0.14218309453938663  
mean of count for weekday 5 is = 0.1431037459002244  
mean of count for weekday 6 is = 0.2885091202025433
```



## 1.4 cleaning up data

```
[4]: data = d.drop(['instant', 'casual', 'registered'], axis=1)
# casual and registered are not features : they are certain predictors of count
# probably meant to be used in the anomaly detection challenge
# sticking to zero-indexed data
```

## 1.5 splitting data:

using a phase indicator: train as 0, val as 1 and test as 2

```
[5]: data['split'] = 0
```

```
[6]: # data.loc[data['dteday'].str[-2:].astype(int) < 20, 'split'] = 0 #train split
data.loc[data['dteday'].str[-2:].astype(int) >= 20, 'split'] = 2 # test split
assert len(data) == data[data['split']==0]['split'].
    ↳count()+data[data['split']==2]['split'].count()
```

```
[7]: # ## validation split
# choosing 20 percent of the train data to be validation
# starting after the first 100 records as validating right away makes no sense
val_candidates = [i for i in range(100, len(data)) if data.iloc[i]['split']==0 ]
val_chosen = random.sample(val_candidates, int(0.2*(len(val_candidates) +100)))
```



```
data.iloc[val_chosen,-1] = 1
```

```
[8]: # ## some final processing
data = data.drop('dteday',axis=1) # not needed anymore
```

```
[9]: norm = ['weathersit','season','mnth','hr','weekday']
# ones to be normalized
# could check a one hot-encoded version of season as well
# but they follow an order so going with this for now
for n in norm:
    data.loc[:,n] = data[n]/max(data[n])
```

main routines have been well documented in the model.py module in the same directory proceeding with analysis in the notebook from now on

## 1.6 this is the report with experiments

## 1.7 check model.py for code : it is self-explanatory(well-commented)

## 2 Poisson regression

- modelling the rate per hour to be parameter of a poisson model
- as the count is always positive, modelling the parameter as a log linear model

hence for the parameters  $W_{d-vect}$  and features  $X_{d-vect}$ , the probability density is given as:

$$f(y) = \frac{\lambda^y \cdot e^{-\lambda}}{y!}$$

where

$$\lambda = e^{W^T X}$$

### 2.1 some nuances

- we will have to use an iterative method for optimization as a closed form solution won't be possible, due to the exponential term in the log-likelihood
- using the negative-log likelihood as the loss function
- using an lr  $\eta$
- for L1 and L2 regression, introducing hyper-parameters  $\alpha$  and  $\beta$  respectively and these will be altered in the validation phase via grid search
- only the previous data can be used for prediction hence we will have to update the weight vector intermittently
- proceeding with a batch size of 1
- will have modes for processing : train, val and test
- will differentiate w.r.t  $W$  in the train mode
- will differentiate w.r.t hyper-parameter in val mode
- will only output prediction in test mode
- normalizing some features to the train set and fitting the test to it while prediction, for better performance and also for the last sub-question

### 3 subquestion 1

#### 4 Likelihood and Loss function

for an  $X, y$  pair: the likelihood will be

$$\mathcal{L}(W) = \frac{e^{-\lambda} \lambda^y}{y!}$$

$$\therefore \ln(\mathcal{L}(W)) = y \ln(\lambda) - \lambda - \ln(y!)$$

ignoring the constant factorial from here onwards

$$\ln(\mathcal{L}(W)) = yW^T X - e^{W^T X}$$

employing the loss function to be the negative log likelihood

$$\therefore L(loss) = e^{W^T X} - yW^T X$$

also note the derivative:

$$\nabla_W L = e^{W^T X} X - yX = X(e^{W^T X} - y)$$

#### 4.1 L1 regularization

for  $\alpha$  as hyper-parameter

we have an additional term of  $\alpha \|W\|$  with the loss function

$$\nabla_W L_{L1} = (\text{existing term}) \cdots + \alpha \nabla \|W\|$$

$$\|W\| = (\|W\|^2)^{\frac{1}{2}}$$

$$\therefore \nabla \|W\| = \frac{1}{2} \|W\|^{-\frac{1}{2}} \cdot 2W = \frac{W}{\|W\|}$$

overall(for training phase)

$$\therefore \nabla_W L_{L1} = X(e^{W^T X} - y) + \alpha \frac{W}{\|W\|}$$

#### 4.2 L2 regularization

for  $\beta$  as hyper-parameter

we have an additional term of  $\beta \|W\|^2$  with the loss function

$$\nabla_W L_{L2} = (\text{existing term}) \cdots + \beta \nabla \|W\|^2$$

overall(for training phase)

$$\therefore \nabla_W L_{L2} = X(e^{W^T X} - y) + 2\beta W$$

#### 4.3 Experiments

```
[60]: import importlib
import model
```

```
[61]: importlib.reload(model)
# for quick retesting
```

```
[61]: <module 'model' from '/mnt/c/leisure and imp docs/BTECH CSE 4 yrs/3rd
year/sem5/Foundations of Machine Learning/fml-assignment-1/Question 4/model.py'>
```

```
[103]: base = model.BaseModel(0.0001,13,data,-1,-2);
L1 = model.L1Model(0.0001,13,data,-1,-2,0.0001);
L2 = model.L2Model(0.0001,13,data,-1,-2,0.0001);
```

Notes: - note that the question demands that we use only the previous data to predict what comes next - for this we predict whenever we encounter a test split index right away - also, we are processing the index one by one ( batch-size=1 ) - validating at random times from the train set (except the first 100 train split indices)

check model.py : it has been completely documented

#### 4.4 full pass on dataset

base model

```
[104]: base.full_pass()
```

```
100%|      | 17379/17379 [00:13<00:00, 1327.98it/s]
```

L1 model

```
[105]: L1.full_pass()
```

```
100%|      | 17379/17379 [00:13<00:00, 1298.18it/s]
```

L2 model

```
[106]: L2.full_pass()
```

```
100%|      | 17379/17379 [00:13<00:00, 1327.64it/s]
```

#### 4.5 errors

```
[107]: print('base error:', base.test_rms())
print('L1 error  :', L1.test_rms())
print('L2 error  :', L2.test_rms())
```

```
base error: 214.5731040277971
```

```
L1 error  : 155.5828397399973
```

```
L2 error  : 159.57408948642873
```

## 4.6 note the performance improvement with the same lr with regularization

checking final W:

```
[108]: base.W
```

```
[108]: array([ 3.24786636,  3.04850391, -0.99508023, -0.70369008,  0.52856336,  
          -0.21777155,  0.14183319,  0.03915093, -0.20839004,  1.41613996,  
          0.88802293, -1.11314864,  0.92164374])
```

```
[109]: L1.W
```

```
[109]: array([ 3.05861708,  0.19296796,  0.16451219,  0.97007019,  0.53058912,  
          -0.18821904,  0.18368329,  0.08141973, -0.09701366,  1.11878637,  
          1.41180134, -1.13397959,  1.11854914])
```

```
[110]: L2.W
```

```
[110]: array([ 3.69706063,  0.33118782,  0.0797629 ,  0.10201965,  0.54241072,  
          -0.08303447,  0.24122738,  0.08712386, -0.07443048,  2.16669895,  
          0.40780129, -1.00287435,  1.2805619 ])
```

### 4.6.1 full disclosure:

- the grid search in our model.py is not completely functional and so has been commented out for now

## 4.7 subquestion 5

4.7.1 for checking which features are the most important: looking at the absolute value of the weights

```
[142]: features = list(data.columns[:12])  
features
```

```
[142]: ['season',  
        'yr',  
        'mnth',  
        'hr',  
        'holiday',  
        'weekday',  
        'workingday',  
        'weathersit',  
        'temp',  
        'atemp',  
        'hum',  
        'windspeed']
```

using the weights from the selection operator for now:

```
[129]: weights = L1.W
```

```
[131]: weights  
# 0 is the bias
```

```
[131]: array([ 3.05861708,  0.19296796,  0.16451219,  0.97007019,  0.53058912,  
        -0.18821904,  0.18368329,  0.08141973, -0.09701366,  1.11878637,  
         1.41180134, -1.13397959,  1.11854914])
```

note the large values for index number 3,4,9,10,11,12

these can be used to infer the most prominent features

```
[144]: prominent = [features[i-1] for i in [3,4,9,10,11,12]]  
prominent
```

```
[144]: ['mnth', 'hr', 'temp', 'atemp', 'hum', 'windspeed']
```

Note that these make sense at the first look: - more bikers later in the year than before(for month)  
- temperature: people get out on a sunny day - feeling temperature: as expected - this is a stronger predictor than temperature - hum : less people get out on humid days - windspeed : who doesn't like to ride on windy days...

```
[ ]:
```